

Lecture 1: Introduction and Overview

Information Retrieval
Computer Science Tripos Part II

1 Motivation

- Definition of “Information Retrieval”
- IR: beginnings to now

2 First Boolean Example

- Term-Document Incidence matrix
- Practicalities of Boolean Search

3 Reading

What is Information Retrieval?

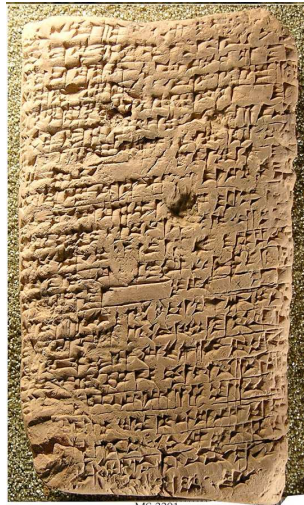
Manning et al, 2008:

Information retrieval (IR) is finding material ... of an unstructured nature ... that satisfies an information need from within large collections

What is Information Retrieval?

Manning et al, 2008:

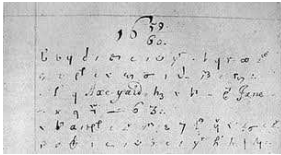
Information retrieval (IR) is finding material ... of an unstructured nature ... that satisfies an information need from within large [collections](#)



MS 3391

Library catalogue. Babylonia, 2000-1600 BC

Document Collections



IR in the 17th century: Samuel Pepys, the famous English diarist, **subject-indexed** his treasured 1000+ books library with key words.

Document Collections

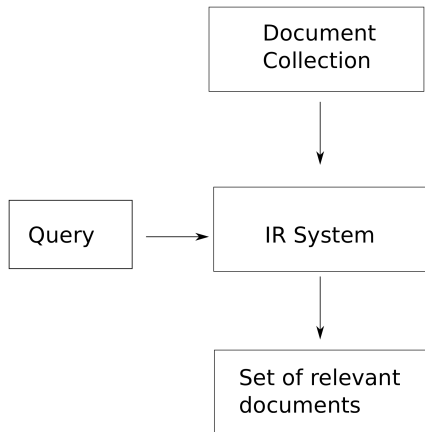



What we mean here by document collections

Manning et al, 2008:

Information retrieval (IR) is finding material (usually documents) of an unstructured nature ... that satisfies an information need from within large **collections (usually stored on computers)**.

- **Document Collection**: text units we have built an IR system over.
- Usually documents
- But could be
 - memos
 - book chapters
 - paragraphs
 - scenes of a movie
 - turns in a conversation...
- Lots of them





What is Information Retrieval?

Manning et al, 2008:

Information retrieval (IR) is finding material (usually documents) of an **unstructured** nature . . . that satisfies an information need from within large collections (usually stored on computers).

Structured vs Unstructured Data

Unstructured data means that a formal, semantically overt, easy-for-computer structure is missing.

- In contrast to the rigidly structured data used in DB style searching (e.g. product inventories, personnel records)



Search Businesses

Name / Type
florists

Location
CB1

[Advanced Business Search](#)

```
SELECT *  
FROM business_catalogue  
WHERE category = 'florist'  
AND city_zip = 'cb1'
```

- This does not mean that there is no structure in the data
 - Document structure (headings, paragraphs, lists...)
 - Explicit markup formatting (e.g. in HTML, XML...)
 - Linguistic structure (latent, hidden)

Manning et al, 2008:

Information retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that **satisfies an information need** from within large collections (usually stored on computers).

- An **information need** is the topic about which the user desires to know more about.
- A **query** is what the user conveys to the computer in an attempt to communicate the information need.
- A document is **relevant** if the user perceives that it contains information of value with respect to their personal information need.

Manning et al, 2008:

Information retrieval (IR) is finding material . . . of an unstructured nature . . . that satisfies an **information need** from within large collections

- Known-item search
- Precise information seeking search
- Open-ended search (“topical search”)

Information scarcity vs. information abundance

- **Information scarcity problem** (or needle-in-haystack problem): hard to find rare information
 - Lord Byron's first words? 3 years old? Long sentence to the nurse in perfect English?

... when a servant had spilled an urn of hot coffee over his legs, he replied to the distressed inquiries of the lady of the house, 'Thank you, madam, the agony is somewhat abated.' [not Lord Byron, but Lord Macaulay]

- **Information abundance problem** (for more clear-cut information needs): redundancy of obvious information
 - What is toxoplasmosis?

Manning et al, 2008:

Information retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that **satisfies** an information need from within large collections (usually stored on computers).






- Are the retrieved documents
 - about the target subject
 - up-to-date?
 - from a trusted source?
 - satisfying the user's needs?
- How should we rank documents in terms of these factors?
- More on this in a lecture soon

How well has the system performed?

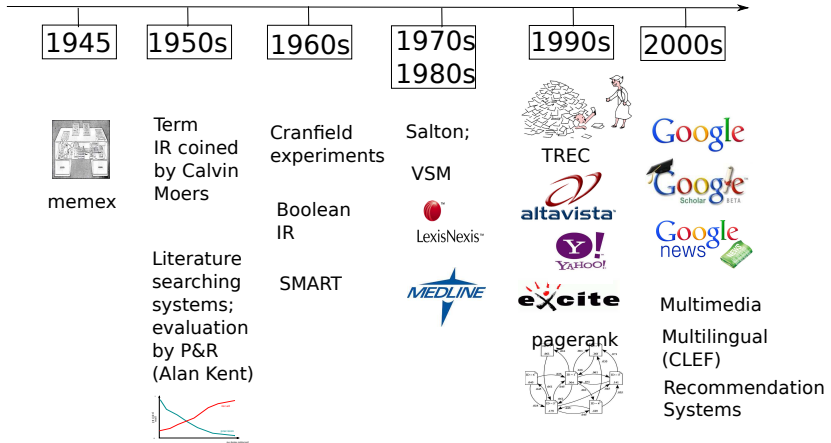
The **effectiveness** of an IR system (i.e., the quality of its search results) is determined by two key statistics about the system's returned results for a query:

- **Precision:** What fraction of the returned results are relevant to the information need?
- **Recall:** What fraction of the relevant documents in the collection were returned by the system?
- What is the best balance between the two?
 - Easy to get perfect recall: just retrieve everything
 - Easy to get good precision: retrieve only the most relevant

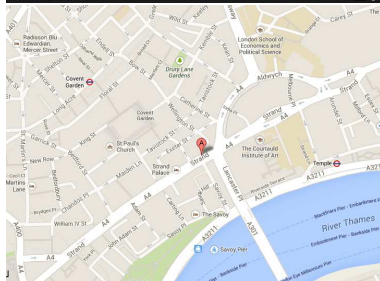

There is much more to say about this – lecture 5

- Web search ( )
 - Search ground are billions of documents on millions of computers
 - issues: spidering; efficient indexing and search; malicious manipulation to boost search engine rankings
 - Link analysis covered in Lecture 8
- Enterprise and institutional search ( )
 - e.g company's documentation, patents, research articles
 - often domain-specific
 - Centralised storage; dedicated machines for search.
 - Most prevalent IR evaluation scenario: US intelligence analyst's searches
- Personal information retrieval (email, pers. documents; )
 - e.g., Mac OS X Spotlight; Windows' Instant Search
 - Issues: different file types; maintenance-free, lightweight to run in background


A short history of IR





IR for non-textual media



Similarity Searches


Reverse Image Search




JPEG, 600x364, 84.9 KB

4 Results
Searched over **4.560 billion** images in 0.837 seconds
for file: trial.jpg

- These results expire in 72 hours. [View](#)
- [Share a success story](#)
- TinEye is [free](#) to use for non-commercial purposes.

Sort by:

Best Match
Most Changed
Biggest Image
Newest
Oldest



Same file | [Link](#)
JPEG image
600x364, 84.9 KB

trendland.com

[kitchen-portraits-by-erik-klein-wolfe...](#)
[trendland.com/kitchen-portraits-by-er...](#)
Crawled on 2013-09-30



[Compare](#) | [Link](#)
JPEG image
600x364, 84.4 KB

freshome.com

[kitchen-portraits-4.jpg](#)
[freshome.com/2013/06/17/kitchen-as-me...](#)
Crawled on 2013-02-12
[freshome.com/2013/06/17/kitchen-as-me...](#)
Crawled on 2014-01-15



[Compare](#) | [Link](#)
JPEG image
640x385, 139.4 KB

trendland.com

[kitchen-portraits-by-erik-klein-wolfe...](#)
[trendland.com/kitchen-portraits-by-er...](#)
Crawled on 2013-09-30



[Compare](#) | [Link](#)
JPEG image
790x470, 101.5 KB

bigpicture.ru

[KitchenPortraits02.jpg](#)
[bigpicture.ru/?p=292266](#)
Crawled on 2013-11-30

GET SHAZAM





SHAZAM MUSIC

TAG CHART FIND MUSIC BLOG INTERVIEWS SHAZAM

Tag Chart - World
The top tracks tagged by Shazamers worldwide, week ending January 05 2014

Track samples provided courtesy of iTunes

World

- **Counting Stars**
OneRepublic
- **Let Her Go**
Passenger
- **Timber**
Pitbull Feat. Ke\$ha
- **Say Something**

- “Ad hoc” retrieval ([lectures 1-5](#))
- web retrieval ([lecture 8](#))
- Support for browsing and filtering document collections:
 - Clustering ([lecture 6](#))
 - Classification; using fixed labels (common information needs, age groups, topics; [lecture 7](#))
- Further processing a set of retrieved documents, e.g., by using natural language processing
 - Information extraction
 - Summarisation
 - Question answering

1 Motivation

- Definition of “Information Retrieval”
- IR: beginnings to now


2 First Boolean Example

- Term-Document Incidence matrix
- Practicalities of Boolean Search

3 Reading

Boolean Retrieval

- In the Boolean retrieval model we can pose any query in the form of a Boolean expression of terms
- i.e., one in which terms are combined with the operators **and**, **or**, and **not**.
- Shakespeare example



Brutus AND Caesar AND NOT Calpurnia

- Which plays of Shakespeare contain the words **Brutus** and **Caesar**, but not **Calpurnia**?
- Naive solution: linear scan through all text – “grepping”
- In this case, works OK (Shakespeare’s Collected works has less than 1M words).
- But in the general case, with much larger text collections, we need to **index**.
- Indexing is an offline operation that collects data about which words occur in a text, so that at search time you only have to access the precompiled index.

The term-document incidence matrix

Main idea: record for each document whether it contains each word out of all the different words Shakespeare used (about 32K).

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0
...						

Matrix element (t, d) is 1 if the play in column d contains the word in row t , 0 otherwise.

Query “Brutus AND Caesar AND NOT Calpurnia”

We compute the results for our query as the bitwise AND between vectors for Brutus, Caesar and complement (Calpurnia):

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0
...						

Query “Brutus AND Caesar AND NOT Calpurnia”

We compute the results for our query as the bitwise AND between vectors for Brutus, Caesar and complement (Calpurnia):

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
¬Calpurnia	1	0	1	1	1	1
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0
...						

Query “Brutus AND Caesar AND NOT Calpurnia”

We compute the results for our query as the bitwise AND between vectors for Brutus, Caesar and complement (Calpurnia):

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
¬Calpurnia	1	0	1	1	1	1
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0
AND						

Query “Brutus AND Caesar AND NOT Calpurnia”

We compute the results for our query as the bitwise AND between vectors for Brutus, Caesar and complement (Calpurnia):

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
¬Calpurnia	1	0	1	1	1	1
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0
AND	1	0	0	1	0	0

Bitwise AND returns two documents, “Antony and Cleopatra” and “Hamlet”.

The results: two documents

Antony and Cleopatra, Act III, Scene ii

Agrippa [Aside to Domitius Enobarbus]: Why, Enobarbus,
When Antony found Julius Caesar dead,
He cried almost to roaring, and he wept
When at Philippi he found Brutus slain.

Hamlet, Act III, Scene ii

Lord Polonius: I did enact Julius Caesar: I was killed i' the
Capitol; Brutus killed me.

- Provided by large commercial information providers
1960s-1990s
- Complex query language; complex and long queries
- Extended Boolean retrieval models with additional operators – proximity operators
- Proximity operator: two terms must occur close together in a document (in terms of certain number of words, or within sentence or paragraph)
- Unordered results...

Example: Westlaw

- Largest commercial legal search service – 500K subscribers
- Boolean Search and ranked retrieval both offered
- Document ranking only wrt chronological order
- Expert queries are carefully defined and incrementally developed

Westlaw Queries/Information Needs

“trade secret” /s disclos! /s prevent /s employe!

Information need: Information on the legal theories involved in preventing the disclosure of trade secrets by employees formerly employed by a competing company.

disab! /p access! /s work-site work-place (employment /3 place)

Information need: Requirements for disabled people to be able to access a workplace.

host! /p (responsib! liab!) /p (intoxicat! drunk!) /p guest

Information need: Cases about a host's responsibility for drunk guests.

- Proximity operators: /3= within 3 words, /s=within same sentence /p =within a paragraph
- Space is disjunction, not conjunction (This was standard in search pre-Google.)
- Long, precise queries: incrementally developed, unlike web search
- Why professional searchers like Boolean queries: precision, transparency, control.

Does Google use the Boolean Model?

On Google, the default interpretation of a query $[w_1 w_2 \dots w_n]$ is $w_1 \text{ AND } w_2 \text{ AND } \dots \text{ AND } w_n$

- Cases where you get hits which don't contain one of the w_i :
 - Page contains variant of w_i (morphology, misspelling, synonym)
 - long query (n is large)
 - Boolean expression generates very few hits
 - w_i was in the *anchor text*
- Google also *rank*s the result set
 - Simple Boolean Retrieval returns matching documents in no particular order.
 - Google (and most well-designed Boolean engines) rank hits according to some estimator of relevance

1 Motivation

- Definition of “Information Retrieval”
- IR: beginnings to now

2 First Boolean Example

- Term-Document Incidence matrix
- Practicalities of Boolean Search

3 Reading

- Manning, Raghavan, Schütze: Introduction to Information Retrieval (MRS), chapter 1
- MRS chapter 2.3